

Fake News Detection Using Machine Learning Algorithms

Akshat Wadhwa
IIIT Delhi
New Delhi, India

akshat19231@iiitd.ac.in

Shruti Jha
IIIT Delhi
New Delhi, India

shruti19274@iiitd.ac.in

Tarini Sharma
IIIT Delhi
New Delhi, India

tarini19451@iiitd.ac.in

1. Introduction

With the rise in popularity of the Internet, social media platforms like Facebook and Twitter have become new sources of information for people. But these online media platforms are being manipulated by certain entities to spread distorted facts. This phenomenon is known as fake news [6]. Fake news can lead to serious problems like the spread of misinformation about vaccines during the COVID-19 pandemic that led to issues in maintaining public health [5]. In this project, we aim to use machine learning algorithms for automated classification of news articles as fake or real. We aim to explore various textual properties in natural language processing on the dataset, which we will use to train different ML models and ensemble methods and evaluate their performance to determine the best model for this learning task.

2. Related Work

In [8], the authors performed very minimal preprocessing: lowercasing the dataset and dropping stopwords. In order to train the models, a 5-fold split was also performed (80% training, 20% test). The models used are CNN, XG-Boost, SVM, Naive Bayes and Random Forest. After training using the Keras API and TensorFlow module, the Accuracy, Precision, Recall and F1-Score are calculated (Figure 5). In paper [7], in addition to lowercasing and removing stopwords, empty strings and missing labels are also handled. The models used are SVM, CNN and LSTM on the metrics Accuracy, Precision and Recall (Figure 6). In [9], the preprocessing is done using word2vec (after dropping the stopwords) from the gensim library. The models used are Feedforward Neural Network, CNN with one convolution layer, CNN with multiple convolution layers and LSTM on the metrics Accuracy, Precision, Recall and F1-score. The metrics are calculated on two different aspects of this dataset: in the first part, the models are trained using the full text of each article and in the latter part models are trained on the title of each article (Figures 7 and 8).

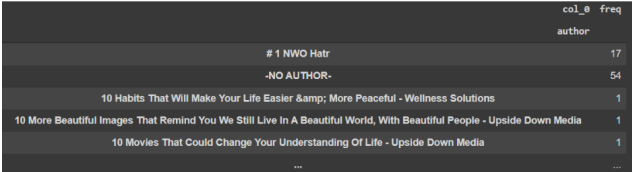
3. Dataset and Evaluation

3.1. Dataset Information

We have used the dataset available on Kaggle [1], which has two files, train.csv (20387 training samples) and test.csv (5127 testing samples). The testing data has four attributes: id, title, author, text and the training data has the additional column of class label (0 for reliable/real news and 1 for unreliable/fake news). There is a slight imbalance in the class distribution in training data, with 10361 samples of class 0 and 7850 samples of class 1 (see Figure 9).

3.2. Pre-processing and feature extraction

Before extracting features, we remove the NaN values and data samples containing only white spaces or “\n”. The entire text is converted to lowercase and English letters from a-z are kept, everything else replaced with space. Then stemming is performed to reduce the words in the dataset to their basic roots. Next we drop stop words like “an”, “the” etc. [2] since they do not carry any significant meaning. Lastly, we vectorize the document by sklearn’s TF-IDF Vectorizer [10]. We explored three different vectorizers for feature extraction before choosing TF-IDF because it gave best results on validation set, as seen in Figure 10. Author attribute had 1957/20800 NaN values in the train set and 500/5200 NaN values in test set. Instead of removing these, replaced NaN values with ‘-no author-’ which helped increase the dataset size. After preprocessing, we



col_0	freq
# 1 NWO Hstr	17
-NO AUTHOR-	54
10 Habits That Will Make Your Life Easier & More Peaceful - Wellness Solutions	1
10 More Beautiful Images That Remind You We Still Live In A Beautiful World, With Beautiful People - Upside Down Media	1
10 Movies That Could Change Your Understanding Of Life - Upside Down Media	1
...	...

Figure 1. Entries with no-author in the author field

had 5040 test samples and 20126 training samples, upon which stratified 5 fold cross validation was performed with 4 folds[16100 samples] used to train the model and 1 fold

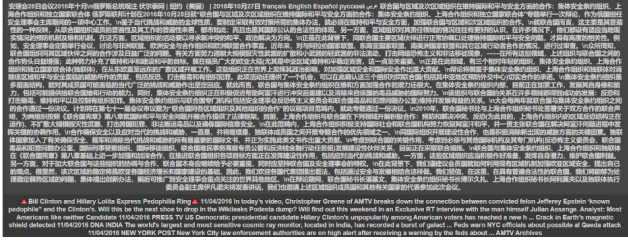


Figure 2. Texts with Chinese and special characters

kept for validation[4026 samples]. Some texts are not in English (Chinese, or have some symbols).

3.3. Evaluation Metrics

Evaluation metrics used are accuracy, recall, precision and F1 score. We also plot the confusion matrix and AUC-ROC curves of the models to better compare them. To evaluate models, a single-number evaluation metric is needed. Since F1-score is the harmonic mean of precision and recall, we choose F1 score as a satisfying metric(threshold=0.6) and accuracy as the optimizing metric. Formulas to determine these metrics are in Figure 11.

4. Analysis and Progress

4.1. EDA

We checked the separability of training data by reducing the dimensions of data to 2 using t-SNE and plotting the scatterplot (Figure 12). We see that the data is not very well separated and there are quite a few samples of class 0 overlapping with samples of class 1. We also plotted the word clouds for the reliable news articles and the fake news articles as shown in Figures 13 and 14. There are some general words like “one”, “time”, “many” present in both real and fake news in equal weightage. To decide the number of features to extract in the tf-idf vectorizer, we plotted the mean dev set (for 5-fold) accuracy with Naive Bayes model and got the best result for 3000 features (Figure 15).

4.2. Models and Learning methods

We have explored 6 ML models: Naives bayes (baseline), Logistic Regression [learning methods used: LBFGS (variant of Newton’s method using approximated Hessian matrix) and SGD], Passive-Aggressive Classifier, SVM [learning methods used: Sequential Minimal Optimisation (SMO) and SGD], Decision Tree and Multi-layer Perceptron [learning methods used: SGD and Quasi-Newton method]. We also explored an RNN based (LSTM) model. [3] We also explored two ensemble methods: XGBoost (boosting method) and Random Forest (bagging method). We used stratified k-fold cross validation (k=5) while training all models and selected the model with best dev set ac-

curacy while satisfying threshold of 0.9 on f1 score.

4.2.1 Hyperparameter Tuning

For hyperparameter tuning, we used Bayesian Optimization implemented in sklearn to get the best hyperparameters for each model (see Figure 16). Since we did not have enough RAM to fit the whole dataset, we used PCA decomposition to reduce the data dimensionality to 100 dimensions.

4.3. Challenges and their Resolutions

Overfitting was detected using learning curves. (see fig. 17-24) Naive Bayes has a high bias of about 10% and a low variance of about 2%. The model is underfitting and so we use advanced models. Logistic regression[LBFGS] has a low bias of 0.01%(assuming desired performance to be 100% accuracy) and has a variance of 6%. Logistic regression[SGD] has a bias of 2% and variance of 3.4%. Passive Aggressive model has bias of 1.3% and variance of 4.7%. Random Forest model has a bias of 0.01% and variance of about 5%. XGBoost has 2% bias and 3% variance. All the models except Naive Bayes is fitting well to the data. To reduce overfitting while training, we tried to add regularization(L2 regularization for logistic regression, early stopping in SGD). We also tried decreasing the number of features which decreased the dev set accuracy as shown in EDA). Reducing the number of features to 100 increased bias by 14.0%

4.4. Covariate shift check between validation set and test set

Preprocessed the train(and test) set, added a column with targetDataset=0(and 1) and dropped the classLabel from train dataset. Concatenated and shuffled the 2 datasets and used a RF classifier to classify combined dataset based on targetDataset. An ROC value of 0.505 was observed which implies that the datasets do not have much covariate shift.

4.5. Sentiment Analysis

Sentiment analysis of ‘text’ attribute using vader and textblob. Using textblob, extracted features maxpos, minneg, mean_sentiment, mean_subjectivity. Textblob.sentiment.polarity returns polarity of text in the range [-1,1] where -1 denotes negative sentiment and +1 denotes positive sentiment. Textblob.sentiment.subjectivity returns the amount of personal opinion and factual information in the range [0,1] where 1 denotes that the text contains more personal opinion rather than factual information. [4] Maxpos is the maximum polarity over all sentences in the ‘text’ attribute of a single datasample. Minneg denotes the minimum polarity, mean_sentiment denotes the mean polarity. Mean_subjectivity is the mean of subjectivity over all sentences in the datasample ‘text’ attribute. Trained

Model	Accuracy	Precision	Recall	f1-score
Logistic regression [LBFGS]	0.94051704415	0.93260204244	0.9291401273	0.9308614855
Logistic Regression [SGD]	0.94893959909	0.94163003663	0.939826959	0.940722942
Passive-Aggressive	0.94103485169	0.93893152398	0.923393940	0.931008930

Figure 3. Metrics on the validation set

Model	Accuracy	Precision	Recall	f1-score
Logistic regression [LBFGS]	0.94043370848	0.93961038961	0.9210693825	0.9302475088
Logistic Regression [SGD]	0.95004117485	0.94095238095	0.943348185	0.942148760
Passive-Aggressive	0.94125720559	0.94433529796	0.917886696	0.930923176

Figure 4. Metrics on the test set

logistic regression(LBFGS) on the extracted features which results in the following metrics - Accuracy: 0.63107, Precision: 0.601618, Recall: 0.426114, F1-score: 0.498881431. Using vader, extracted features pos, neg, neu, compound, where pos denotes positive sentiment score, neg denotes negative sentiment score and neu denotes neutral sentiment score in the range [0,1] (all sum upto 1). Compound denotes the compound score calculated as the sum of normalized sentiment scores. Trained logistic regression(LBFGS) on the extracted features which results in Accuracy: 0.6184, Precision: 0.6898, Recall: 0.20828, F1-score: 0.31996.

Because of better performance of textblob, concatenated the 4 features from textblob with the 3000 features from TF-IDF and trained on Logistic Regression(LBFGS and SGD) and Passive Aggressive Classifier (shown in figure below).

5. Results and Interpretation

5.1. Results

Figures 45 and 46 contain the value of the metrics mentioned in the section 3 for the validation and test set. Figure 47 contains the accuracy of the models on the test set provided by the Kaggle competition from which the dataset was taken. [1] Gaussian Naive Bayes is underfitting as can be inferred from the learning curve and relatively low training, dev and test accuracy for Naive Bayes. Since we needed to submit our predictions on kaggle to get the accuracy, we decided to do a 80-20 train test split to test our models on other metrics (precision, recall, f1-score, roc curve, confusion matrix) With weighted accuracy as the optimizing metric and F1-score as the satisficing metric(with threshold 0.9), the best model is SVM(SMO) with a test accuracy of about 0.9548

5.2. Interpretation

Looking at dev set samples which are mislabelled as real, it was noticed that the dev set contains some key words(essential for a human classifying the text as fake or real) not present in the train set. Eg, in the validation sample with text “pokemon go players are inadvertently stopping people committing suicide in japan“, “pokemon” and “inadvertently” are not present in the features extracted using TF-IDF vectorizer. This implies we need to make the dataset more diverse in the future. Performance of gaussian naive bayes on the test set is better as comparable to paper [6] (difference in accuracy about 28.8%). Performance of SVM on test set is better than the papers [5] and [6] with difference in accuracy of about 35% and 32%. Performance of random forest and xgboost is better as compared to [6] (difference of about 16% and 6% respectively). This can be attributed to the fact that the authors performed very minimal preprocessing - lowercasing the whole dataset and dropping stopwords. Performance of LSTM is comparable to [8] (difference in performance of)

On comparing the AUC-ROC curves (see figures 25-34) for the various baseline (Gaussian Naive Bayes) and advanced models implemented, we see that the the area under the ROC curve for baseline model is less as compared to the the area of the other advanced models like logistic regression, SVM, etc., and the ensemble methods explored like XG-Boost and Random Forest. This shows that the advanced models are performing better on the test set as compared to the baseline model, as we would have expected.

Consider class 0 (real news) as the negative class and class 1 (fake news) as the positive class. First analyzing the confusion matrix of our baseline model (Gaussian Naive Bayes), we get that the number of true positives and true negatives are almost equal, as well as the number of false positives and false negatives are also almost equal. Also the total number of samples that it is misclassifying is also very high. This shows that the baseline model is not performing very well on the test set, since it is equally classifying and misclassifying the test data samples (figure 35). For the other advanced models: Logistic regression, Passive Aggressive Classifier, SVM, etc., we can see that the number of true positives and true negatives are quite high, while the number of false positives and false negatives are relatively lower (see figures 36-44). Thus these advanced models give better evaluation metric scores on the test set as well.

6. Contributions

6.1. Deliverables

During the course of the project, the following deliverables were promised and delivered by the corresponding team members:

Shruti - Promised deliverable: EDA, preprocessing, training+hyperparameter tuning of models (SVM, Decision Trees, MLP), model selection, ensemble approaches, drawing final conclusions+report writing.

Delivered deliverables: EDA, preprocessing, training+hyperparameter tuning of models (SVM, Decision Trees, MLP), model selection, ensemble approaches, drawing final conclusions+report writing.

Tarini - Promised deliverables: training+hyperparameter tuning of models (Logistic regression[LBFGS, SGD], Passive Aggressive Classifier, (XGBoost) ensemble approach), model selection, conclusion+report writing.

Delivered deliverables: EDA, training+hyperparameter tuning of models (Logistic regression[LBFGS, SGD], Passive Aggressive Classifier, XGBoost, Random Forest (ensemble approaches)), error analysis, covariate shift, sentiment analysis, model selection, conclusion+report writing.

Akshat - Promised deliverables: Preprocessing, training+hyperparameter tuning of models (Naive Bayes, LSTM), model selection, evaluation metrics, error analysis, drawing final conclusions+report writing.

Delivered deliverables: Pre-processing, Vectorization of data, training+hyperparameter tuning of models (Gaussian Naive Bayes and LSTM), evaluation metrics, drawing final conclusions+report writing.

6.2. File links

During the course of the project, the following files were made by the corresponding team members:

Tarini :

- **Models(training + Hyperparameter tuning)**
- **Learning curves**
- **Error analysis**
- **Preprocessed train and test datasets**

Akshat :

- **TFIDF + Hyperparam tuning**
- **LSTM**
- **Count Vectorizer**
- **Word2vec**

Shruti :

- **EDA**
- **Hyperparameter Tuning**
- **Models (SVM, Decision Trees, MLP)**
- **Learning methods explored for models**
- **Learning curves**

Classifier	Accuracy	Precision	Recall	F1-score
CNN (default)	0.8889	0.8696	0.9006	0.8846
CNN (boosted)	0.9213	0.9150	0.9248	0.9194
XGBoost	0.8992	0.8778	0.9135	0.8953
SVM	0.6311	0.6096	0.6276	0.6184
Naive Bayes	0.5810	0.4801	0.5893	0.5291
Random Forest	0.7853	0.7112	0.8269	0.7647
CNN + XGBoost	0.9487	0.9941	0.9117	0.9511
CNN + SVM	0.8328	0.9736	0.7603	0.8538
CNN + Bayes	0.7921	0.9565	0.7205	0.8219
CNN + Random Forest	0.9458	0.9941	0.9068	0.9485

Figure 5. Evaluation Metrics results in referenced paper [8]

Dataset	Method	Accuracy	Precision	Recall
Liar dataset	SVM	0.608	0.603	0.608
	CNN	0.614	0.611	0.614
	CNN-LSTM	0.623	0.620	0.623
News Articles dataset	SVM	0.683	0.680	0.683
	CNN	0.708	0.701	0.708
	CNN-LSTM	0.725	0.721	0.725

Figure 6. Evaluation Metrics results in the referenced paper [7]

RESULTS ON THE FULL TEXT DATA					
	Accuracy	Precision	Recall	F1	Time
FF	0.898121	0.909406	0.884449	0.896754	26 s
CNN1	0.961705	0.948131	0.976890	0.962295	34 s
CNN2	0.975193	0.969480	0.981178	0.975294	82 s
LSTM	0.918593	0.908197	0.931764	0.919829	906 s

Figure 7. Metrics of the models trained only on text field as given in the referenced paper [9]

RESULTS ON THE TITLE TEXTS					
	Accuracy	Precision	Recall	F1	Time
FF	0.912796	0.854348	0.990923	0.917581	8 s
CNN1	0.911808	0.851036	0.993949	0.916957	4 s
CNN2	0.933547	0.909135	0.958716	0.933267	3 s
LSTM	0.912549	0.846620	1.000000	0.916940	25 s

Figure 8. Metrics of the models trained on the title field as given in the referenced paper [9]

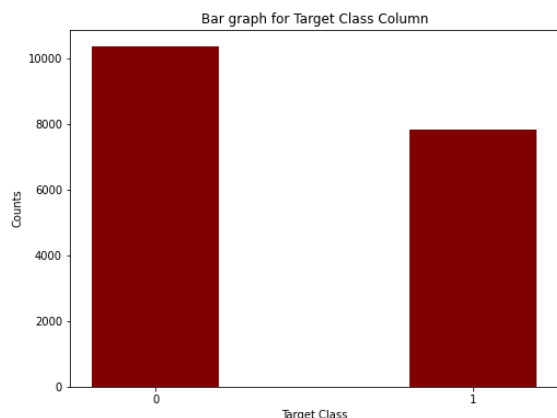


Figure 9. Bar Graph for Target class distribution (Class 0 is real news and Class 1 is fake news)

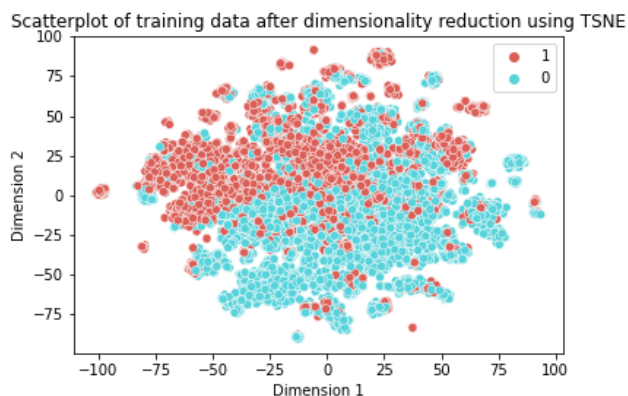


Figure 12. Scatterplot of the training data using t-SNE to check separability of data.

Evaluation metrics on Naive Bayes:

1) Doc2Vec:

Accuracy: 0.7219823232266134

```
f1-score: 0.71041189131552
```

2) CountVectorizer:

Accuracy: 0.8355942044977143

```
f1-score: 0.8277077440158903
```

3) TF-IDF Vectorizer

Accuracy: 0.8818294448984256

```
f1-score: 0.8628961890816955
```

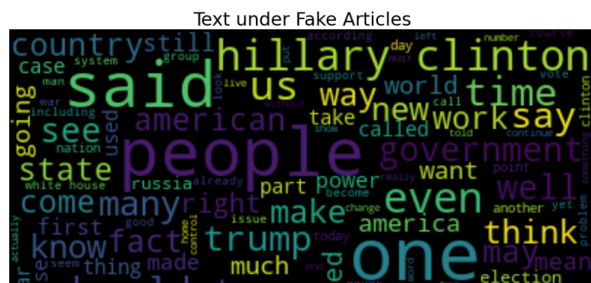


Figure 13. Word cloud for text of Fake articles

Figure 10. Validation accuracy using Baseline model of the three different vectorizers explored for feature extraction.




Figure 11. Formulae for the evaluation metrics used by us.




Figure 14. Word cloud for text of Real articles




Figure 15. Number of features(in tf-idf vectorization of training data) Vs Mean Accuracy of Naive Bayes model(on validation set)

- Naive Bayes (Baseline)
 - var_smoothing - 4.64e-07
- Logistic Regression
 - [LBFGS] : C - 838 (weak regularization), penalty - l2
 - [SGD] : early_stopping - True, penalty - l2, alpha - 1e-05 (weak regularization)
- Passive Aggressive
 - C - 654 (weak regularization)s, early_stopping - True
- SVM
 - [SMO] : C - 3 (strong regularization), kernel - rbf
 - [SGD] : alpha - 0 (weak regularization), learning_rate - constant, penalty - elasticnet
- Decision Tree
 - Criterion - entropy, max_depth - 10, min_samples_leaf - 5, ccp_alpha - 0.0
- Multilayer Perceptron
 - [SGD]: activation - relu, alpha - 1.3171 (strong regularization), learning_rate - constant

Figure 16. Best Hyperparameters obtained using Bayesian Optimisation Hyperparameter tuning of models



Figure 18. Learning Curve for Logistic Regression




Figure 19. Learning Curve for Logistic regression [SGD]




Figure 17. Learning Curve for Naive Bayes




Figure 20. Learning Curve for the Passive Aggressive model




Figure 21. Learning Curve for the Random Forest model




Figure 24. Learning Curve for the Decision Tree model




Figure 22. Learning Curve for the XGBoost model

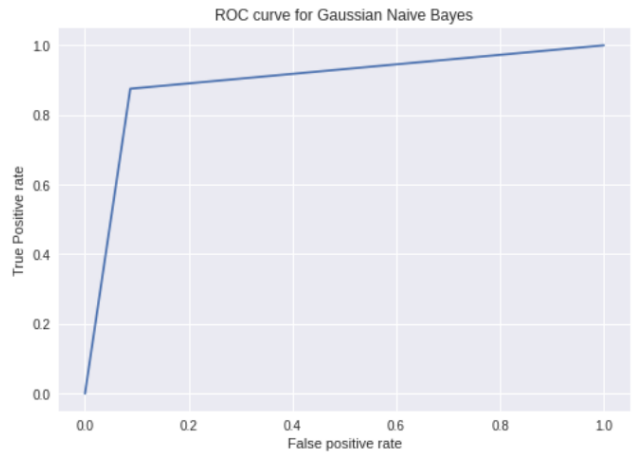


Figure 25. ROC curve for Gaussian Naive Bayes




Figure 23. Learning Curve for the SVM model




Figure 26. ROC curve for Random Forest




Figure 27. ROC curve for XGBoost




Figure 30. ROC curve for Logistic Regression [LBFGS]




Figure 28. ROC curve for Passive Aggressive Classifier




Figure 31. ROC curve for SVM




Figure 29. ROC curve for Logistic Regression [SGD]

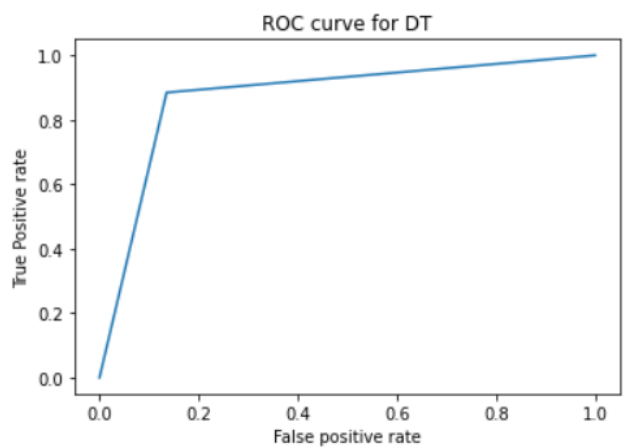


Figure 32. ROC curve for Decision Tree




Figure 33. ROC curve for MLP




Figure 36. Confusion matrix for Random Forest




Figure 34. ROC curve for LSTM




Figure 37. Confusion matrix for XGBoost




Figure 35. Confusion matrix for Gaussian Naive Bayes




Figure 38. Confusion matrix for Passive Aggressive Classifier




Figure 39. Confusion matrix for Logistic Regression [SGD]




Figure 42. Confusion matrix for Decision Tree




Figure 40. Confusion matrix for Logistic Regression [LBFGS]




Figure 43. Confusion matrix for MLP




Figure 41. Confusion matrix for SVM




Figure 44. Confusion matrix for LSTM

Model	Accuracy	Precision	Recall	f1-score
Naives Bayes (baseline)	0.8716086925	0.8779999609	0.853373114	0.86548016
Logistic Regression [LBFGS]	0.94014263092	0.93235222800	0.9284915116	0.930413501
Logistic Regression [SGD]	0.94529099549	0.93675655105	0.936296216	0.936511895
Passive-Aggressive	0.93876964102	0.93257404076	0.924827568	0.928670350
SVM [SMO]	0.9563701923076924	0.950884042674837	0.9622842908705346	0.9565277687363178
SVM [SGD]	0.9492198597	0.94809468932	0.947015943	0.94750381
Decision Tree	0.8808894230769232	0.8639081839218676	0.9036018213904443	0.8832617528890534
Multi-layer Perceptron [SGD]	0.9408653846153847	0.9394999075092343	0.9421612671300028	0.9408135224194648
Multi-layer Perceptron [Quasi-Newton]	0.9265625	0.912003154	0.94384852	0.9276436
Random Forest	0.9441929428	0.9526762669	0.916071256	0.93397766
XGBoost	0.9493409540	0.9385002623	0.944419392	0.94142104

Figure 45. Performance of our models on validation set (5-fold)

Model	Accuracy	Precision	Recall	f1-score
Naives Bayes (baseline)	0.8681073025335321	0.8781535158346753	0.843298969072165	0.8603733894293979
Logistic Regression [LBFGS]	0.94070820752	0.93907971484	0.922342457	0.930635838
Logistic Regression [SGD]	0.95113917101	0.94618834080	0.940165499	0.943167305
Passive-Aggressive	0.94153170463	0.93863049095	0.924888605	0.931708881
SVM [SMO]	0.9548076923	0.9491604477	0.962630085	0.95584781
SVM [SGD]	0.9456736366	0.9411423332	0.954613449	0.94780193
Decision Tree	0.8747596153	0.8709827666	0.884578997	0.87772823
Multi-layer Perceptron [SGD]	0.9423076923	0.9407337723	0.946073793	0.94339622
Multi-layer Perceptron [Quasi-Newton]	0.9230769230	0.91184573002	0.9394512771	0.925442688
Random Forest	0.9475706835	0.9521625163	0.924888605	0.93832741
XGBoost	0.9527861652	0.9385579937	0.952896244	0.94567277
LSTM	0.913065076995032	0.9007056451612904	0.9211340206185566	0.9108053007135576

Figure 46. Performance of our models on the test set.

Model	Accuracy on Kaggle competition
Naives Bayes (baseline)	0.52179
Logistic Regression [LBFGS]	0.93626
Logistic Regression [SGD]	0.94395
Passive-Aggressive	0.94120
SVM [SMO]	0.95549
SVM [SGD]	0.94561
Decision Tree	0.87939
Multi-layer Perceptron [SGD]	0.94285
Multi-layer Perceptron [Quasi-Newton]	0.93049
Random Forest	0.94395
XGBoost	0.94862
LSTM	0.91217

Figure 47. Performance of our models on the test set provided by Kaggle.

Future Tasks	Team Member Assigned
Augmenting training dataset	Tarini, Shruti
Further Error analysis	Everyone
Dropping drift features	Tarini
Adding polarity feature in dataset	Akshat
Explore Doc2vec vectorization	Akshat
Ensemble approaches	Tarini, Shruti
Quasi-Newton Learning technique for MLP model	Shruti
Advanced RNN model (LSTM)	Akshat

Figure 48. Future Tasks distribution amongst the team

References

- [1] [Dataset Available Here \(Click\)](#). 1, 3
- [2] [Pre-processing in NLP \(Click\)](#). 1
- [3] [Referred LSTM's implementation from this link \(Click\)](#). 2
- [4] [TextBlob and Sentiment Analysis referred from here \(Click\)](#). 2
- [5] A. Salma et al. B. Zapan, B. Sajib. Effects of misinformation on covid-19 individual responses and recommendations for resilience of disastrous consequences of misinformation, 2020. Volume 8, ISSN 2590-0617. 1
- [6] Y. Benkler et al. D. M. J. Lazer, M. A. Baum. The science of fake news, 2018. Science, vol. 359, no. 6380, pp. 1094–1096. 1
- [7] Zineb Giordano Silvia. Drif, Ahlem Ferhat Hamida. Fake news detection method based on text-features, 2019. 1, 4
- [8] Anna Marciniak Martyna Tarczewska and Agata Gie lczyk. Fake or real? the novel approach to detecting online disinformation based on multi ml classifiers. 2020. 1, 4
- [9] Viera Maslej Krešňáková, Martin Sarnovsky, and Peter Butka. Deep learning methods for fake news detection. 11 2019. 1, 4
- [10] Ray Oshikawa, Jing Qian, and William Yang Wang. A survey on natural language processing for fake news detection. *arXiv preprint arXiv:1811.00770*, 2018. 1